

MixedComplementarityProblems.jl: A Fast, Batched, Open-Source Interior-Point Solver for Mixed Complementarity Problems

David Fridovich-Keil

Department of Aerospace Engineering and Engineering Mechanics

The University of Texas at Austin

Austin, TX, USA

dfk@utexas.edu

Abstract—Mixed complementarity problems (MCPs) arise as the first-order optimality conditions of nonlinear programs and noncooperative games, and give a natural formulation for the multi-agent trajectory optimization problems that appear throughout robotics. The dominant solver for problems of this form is `PATH`, a mature but closed-source package. We present `MixedComplementarityProblems.jl`, an open-source, pure Julia implementation of an interior-point method for parametric MCPs that (i) matches `PATH`’s reliability on standard benchmarks and (ii) natively supports batched solving, in which many parameter instances are solved simultaneously and in parallel, either across CPU threads or on an NVIDIA GPU, behind a single solver implementation. On a multi-agent lane-change trajectory game representative of robotics planning problems, our CPU-multithreaded batched solver clears a batch of parametric instances 60–80× faster than sequential calls to `PATH`, and our GPU backend clears the same batches roughly 20× faster than `PATH` **[TODO: finalize the GPU-vs-CPU comparison for the trajectory game once the latest benchmark numbers are in—earlier runs showed GPU trailing CPU here, but that may no longer hold]**. We describe the solver’s interior-point formulation, the abstraction that lets a single solver implementation run unmodified across dense, batched-sparse, and single-large linear-algebra backends, and report benchmarks against `PATH` on both randomly generated quadratic programs and trajectory games.

I. INTRODUCTION

Multi-agent trajectory planning problems in robotics, ranging from autonomous driving to human-robot interaction, are naturally posed as noncooperative dynamic games, in which each agent optimizes its own trajectory subject to the others’ choices. The first-order necessary conditions of optimality for each agent in the game take the form of a mixed complementarity problem (MCP). Solving MCPs rapidly, therefore, becomes a core computational burden for real-time decision-making in multi-agent robotics applications.

The dominant solver for this class of problems is `PATH` [1]. `PATH` is mature, reliable, and widely used; but, it is closed-source, solves one problem instance at a time, and was not designed to support automatic differentiation of solutions with respect to problem parameters. All three pose serious, practical limitations for use in robotics: (i) developers must have access to solver internals in order to customize to the problem at hand, (ii) scenario- and contingency-based planning (e.g., [2], [3]) requires solving a *batch* of identically-structured but differently *parameterized* problems, and (iii) modern machine learning pipelines often embed optimization problem solvers as “layers” [4], and end-to-end differentiability requires us to

be able to differentiate a solver’s output with respect to its input parameters.

We present `MixedComplementarityProblems.jl`, an open-source, pure Julia solver for parametric MCPs that addresses all of these practical challenges. Concretely, we contribute:

- A simple and easily-customized interior-point method for parametric MCPs that matches `PATH`’s reliability on standard benchmarks, while supporting automatic differentiation of solutions with respect to problem parameters.
- A native *batched* solve mode, in which many parameter instances sharing one MCP structure are solved simultaneously, parallelized across CPU threads or an NVIDIA GPU behind a single solver implementation.
- An abstraction boundary defined by a small verb set that separates the interior-point loop and its linear algebra backend, which lets a single solver implementation run unmodified across dense, batched-sparse, and single-large-instance regimes, so that supporting a new device or linear-algebra strategy never requires touching the solver loop itself.

As Figure 1 previews and Section V quantifies, on a multi-agent lane-change trajectory game representative of robotics planning problems, clearing a batch of parametric instances with our batched solver takes a small fraction of the wall-clock time required to clear the same batch with sequential calls to `PATH`. **[TODO: replace with the top-line numerical speedup results when they are finalized]**

II. BACKGROUND: MIXED COMPLEMENTARITY PROBLEMS

A. Problem definition

A mixed complementarity problem is specified by a map $F : \mathbb{R}^n \rightarrow \mathbb{R}^n$ and bounds $\ell, u \in (\mathbb{R} \cup \{\pm\infty\})^n$; a solution is a vector $z \in \mathbb{R}^n$ satisfying, for every coordinate dimension i ,

$$\begin{aligned} \ell_i < z_i < u_i &\implies F_i(z) = 0, \\ z_i = \ell_i &\implies F_i(z) \geq 0, \\ z_i = u_i &\implies F_i(z) \leq 0. \end{aligned} \tag{1}$$

Readers are referred to [5] for a complete introduction to MCPs and related problems.

[TODO: replace with rendered lane-change trajectory game schematic, showing the two vehicles' solved trajectories overlaid with a faint fan of trajectories from other initial conditions in the same batch]

(a) One multi-agent trajectory game is one instance out of a batch of parametrically related MCPs, solved simultaneously.

[TODO: replace with wall-clock-vs-batch-size plot: PATH (serial), CPU-batched, and GPU-batched, log-scaled y-axis]

(b) Clearing that batch: sequential PATH calls versus our batched CPU and GPU solves.

Fig. 1. [TODO: finalize caption once the real figure is in] Solving a batch of parametric trajectory-game instances (a) with `MixedComplementarityProblems.jl`'s batched interior-point solver is substantially faster than solving the same batch with sequential calls to `PATH` (b).

B. MCPs from KKT conditions

Consider a parametric nonlinear program

$$\min_x f(x; \theta) \quad \text{s.t.} \quad g(x; \theta) = 0, \quad h(x; \theta) \geq 0, \quad (2)$$

with primal variables $x \in \mathbb{R}^{n_x}$, parameters θ , and Lagrange multipliers $\lambda \in \mathbb{R}^{n_\lambda}$ associated with the equality constraint g and $\mu \in \mathbb{R}_{\geq 0}^{n_\mu}$ associated with the inequality constraint h . The first order necessary (i.e. Karush-Kuhn-Tucker, KKT) conditions of optimality for problem (2) are

$$0 = \begin{bmatrix} \overbrace{\nabla_x f(x; \theta) - \nabla_x g(x; \theta)^\top \lambda - \nabla_x h(x; \theta)^\top \mu}^{G(x, \lambda, \mu; \theta)} \\ g(x; \theta) \end{bmatrix} \quad (3a)$$

$$0 \leq \underbrace{h(x; \theta)}_{H(x; \theta)} \perp \mu \geq 0. \quad (3b)$$

Problem (3) can be converted into the form (1) by taking $z = (x, \lambda, \mu)$, $F = (G, H)$, $u_i = \infty$ ($\forall i$), $\ell_i = -\infty$ ($\forall i \in \{1, 2, \dots, n_x + n_\lambda\}$), and $\ell_i = 0$ ($\forall i > n_x + n_\lambda$). Because of the natural connection between KKT conditions (e.g., arising from trajectory optimization problems in robotics) and this *primal-dual* form of the MCP, it is the only form `MixedComplementarityProblems.jl` exposes.

C. MCPs from noncooperative games

The construction in Section II-B extends from one decision maker to many. Consider N agents, where agent i solves

$$\forall i \begin{cases} \min_{x_i} f_i(x_i, x_{-i}; \theta) \\ \text{s.t.} \quad g_i(x_i, x_{-i}; \theta) = 0, \quad h_i(x_i, x_{-i}; \theta) \geq 0 \end{cases} \quad (4a)$$

$$\text{s.t.} \quad g_{\text{sh}}(x; \theta) = 0, \quad h_{\text{sh}}(x; \theta) \geq 0, \quad (4b)$$

subject to its own equality and inequality constraints (g_i and h_i , respectively) and to constraints $g_{\text{sh}}, h_{\text{sh}}$ shared across

agents (e.g., pairwise collision avoidance). Each agent i 's decision variable is $x_i \in \mathbb{R}^{n_{x_i}}$, and x_{-i} denotes the other agents' decisions so that $x = (x_i, x_{-i}) = (x_1, \dots, x_N)$. Writing out each agent's KKT conditions as in (3) and concatenating them—with the shared constraints $g_{\text{sh}}, h_{\text{sh}}$ coupling agents through shared multipliers $\lambda_{\text{sh}}, \mu_{\text{sh}}$ —yields a single primal-dual MCP whose solution is a generalized Nash equilibrium: a strategy profile from which no agent can unilaterally improve its own objective [5].

Concretely, concatenate the Lagrange multipliers corresponding to equality constraints as $\lambda := (\lambda_1, \dots, \lambda_N, \lambda_{\text{sh}})$, and those for inequality constraints as $\mu := (\mu_1, \dots, \mu_N, \mu_{\text{sh}})$. Then, private multipliers as $\lambda_{\text{pr}} = (\lambda_1, \dots, \lambda_N)$, $\mu_{\text{priv}} = (\mu_1, \dots, \mu_N)$. Each agent i 's Lagrangian is given by

$$\mathcal{L}_i(x, \lambda, \mu; \theta) := f_i(x; \theta) - \lambda_i^\top g_i(x; \theta) - \mu_i^\top h_i(x; \theta) - \lambda_{\text{sh}}^\top g_{\text{sh}}(x; \theta) - \mu_{\text{sh}}^\top h_{\text{sh}}(x; \theta). \quad (5)$$

The N agents' concatenated KKT conditions are then an instance of (3), with

$$\begin{aligned} G &:= (\nabla_{x_1} \mathcal{L}_1, \dots, \nabla_{x_N} \mathcal{L}_N, g_1, \dots, g_N, g_{\text{sh}}), \\ H &:= (h_1, \dots, h_N, h_{\text{sh}}), \end{aligned} \quad (6)$$

with G and H functions of (x, λ, μ) and parameterized by θ .

Running example: A lane-change trajectory game.

Consider the two-agent lane changing trajectory game of Figure 1(a). Here, each agent's decision variable $x_i = \{s_i(t), a_i(t)\}_{t=1}^T$ consists of a state-action *trajectory* over time horizon T . The private equality constraints g_i encode system dynamics $(s_i(t), a_i(t)) \rightarrow s_i(t+1)$, private inequality constraints encode actuator limits $\|a_i(t)\| \leq \bar{a}_i$, and shared inequality constraints encode collision-avoidance $\|P(s_1(t) - s_2(t))\| \geq \Delta$ (with projection

matrix P recovering agents' position). Agents' objective functions f_i encode lane and speed preferences, and penalize control effort. Parameters θ encode the agents' initial states $s_i(1)$, lane preferences, desired speeds, actuator limits $\bar{a}_i$, and collision-avoidance radius Δ .

Observe that the shared collision avoidance constraint—and any nonlinearity in agents' dynamics—renders the feasible set nonconvex.

III. SOLVER DESIGN

`MixedComplementarityProblems.jl` implements an intentionally-simple primal-dual interior point method to solve problem (3). To that end, we introduce a slack variable $z \in \mathbb{R}^{n_\mu}$ and rewrite (3) in terms of $z \geq 0$ and a homotopy parameter $\rho > 0$:

$$K_\rho(x, \lambda, \mu, z; \theta) := \begin{bmatrix} G(x, \lambda, \mu; \theta) \\ H(x; \theta) - z \\ z \odot \mu - \rho \mathbf{1} \end{bmatrix} = 0. \quad (7)$$

Note that the last row of K_ρ involves the elementwise product $z \odot \mu$ and, consequently, the number of equations in K_ρ is equal to the total dimension of the primal, dual, and slack variables (x, λ, μ, z) .

As summarized in Algorithm 1, we proceed in rounds by solving (7) via a regularized Newton method and decaying $\rho \rightarrow 0$ until reaching the user's defined tolerance $\|K_0\| \leq \bar{K}$. For a given value of ρ , the regularized Newton direction $\delta := (\delta_x, \delta_\lambda, \delta_\mu, \delta_z)$ solves

$$\left(\begin{bmatrix} \nabla_x G & \nabla_\lambda G & \nabla_\mu G & 0 \\ \nabla_x H & 0 & 0 & -I \\ 0 & 0 & \text{dg}(z) & \text{dg}(\mu) \end{bmatrix} + \epsilon I \right) \delta = -K_\rho, \quad (8)$$

with $\epsilon \geq 0$ a regularization parameter chosen as described below, and $\text{dg}(\cdot)$ returning a square matrix with its argument on the main diagonal and all other entries zero.

We implement two additional subtleties. First, on line 6, we conduct a “fraction to the boundary” linesearch to control the rate at which $z, \mu \geq 0$ can approach 0 (cf. **[TODO: reference to Nocedal and Wright's book where they talk about this]**). At each iteration k , stepsizes α_k, β_k are chosen so that

$$\begin{aligned} \alpha_k &= \max(\alpha \in [0, 1] : z + \alpha \delta_z \geq (1 - \tau)z) \\ \beta_k &= \max(\beta \in [0, 1] : \mu + \beta \delta_\mu \geq (1 - \tau)\mu), \end{aligned} \quad (9)$$

where the inequalities are interpreted elementwise and $\tau \in (0, 1)$ controls the rate of decay. Linesearch (9) has a closed-form, exact solution, which makes it amenable to batched computation as described in Section IV.

Second, we construct an *adaptive* choice of the Newton step regularization parameter ϵ and the interior point homotopy parameter ρ that aims to control subproblem conditioning. At the end of each inner loop (lines 11 and 13), if the solver successfully drove $\|K_\rho\| \leq \rho$ then we *tighten* ρ and ϵ at a rate dependent upon the number of Newton steps required to converge; conversely, if the inner loop could not meet the tolerance and $\|K_\rho\| > \rho$, we *loosen* both ρ and ϵ to improve the conditioning of the next subproblem and require a less precise solution.

Algorithm 1: Primal-dual interior-point method for (3)

Input: initial iterate (x, λ, μ, z) with $\mu, z > 0$;
parameters θ ; initial $\rho, \epsilon > 0$;
tightening/loosening rates $\underline{\gamma}, \bar{\gamma} \in (0, 1)$; max.
inner iterations $k_{\max}$; tolerance $\bar{K}$

Output: (x, λ, μ, z) with $\|K_0(x, \lambda, \mu, z; \theta)\| \leq \bar{K}$

```

1 while  $\|K_0(x, \lambda, \mu, z; \theta)\| > \bar{K}$  do
2    $k \leftarrow 0$ 
3   repeat
4     assemble  $K_\rho$  (7) and its regularized Jacobian
      (8)
5     solve (8) for the Newton direction
       $\delta = (\delta_x, \delta_\lambda, \delta_\mu, \delta_z)$ 
6     compute step sizes  $\alpha_k, \beta_k$  via the
      fraction-to-boundary rule (9)
7      $x \leftarrow x + \alpha_k \delta_x, \quad \lambda \leftarrow \lambda + \alpha_k \delta_\lambda,$ 
       $z \leftarrow z + \alpha_k \delta_z, \quad \mu \leftarrow \mu + \beta_k \delta_\mu$ 
8      $k \leftarrow k + 1$ 
9   until  $\|K_\rho(x, \lambda, \mu, z; \theta)\| \leq \rho$  or  $k \geq k_{\max}$ 
10  if  $\|K_\rho(x, \lambda, \mu, z; \theta)\| \leq \rho$  then
11     $\rho \leftarrow \rho(1 - e^{-\underline{\gamma}k}), \quad \epsilon \leftarrow \epsilon(1 - e^{-\underline{\gamma}k})$ 
      // tighten
12  else
13     $\rho \leftarrow \rho(1 + e^{-\bar{\gamma}k}), \quad \epsilon \leftarrow \epsilon(1 + e^{-\bar{\gamma}k})$ 
      // loosen
14   $\rho \leftarrow \min(\rho, 1)$ 
15 return  $(x, \lambda, \mu, z)$ 

```

IV. BATCHED AND GPU-PARALLEL SOLVING

[TODO: TODO]

V. EXPERIMENTS

[TODO: TODO]

VI. DISCUSSION AND FUTURE WORK

[TODO: TODO]

REFERENCES

- [1] S. P. Dirkse and M. C. Ferris, “The path solver: a nonmonotone stabilization scheme for mixed complementarity problems,” *Optimization methods and software*, vol. 5, no. 2, pp. 123–156, 1995.
- [2] J. Li, C. Y. Chiu, L. Peters, F. Palafox, M. Karabag, J. Alonso-Mora, S. Sojoudi, C. Tomlin, and D. Fridovich-Keil, “Scenario-game admm: A parallelized scenario-based solver for stochastic noncooperative games,” in *62nd IEEE Conference on Decision and Control, CDC 2023*. IEEE, 2023, pp. 8093–8099.
- [3] L. Peters, A. Bajcsy, C.-Y. Chiu, D. Fridovich-Keil, F. Laine, L. Ferranti, and J. Alonso-Mora, “Contingency games for multi-agent interaction,” *IEEE Robotics and Automation Letters*, vol. 9, no. 3, pp. 2208–2215, 2024.
- [4] B. Amos and J. Z. Kolter, “Optnet: differentiable optimization as a layer in neural networks,” in *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, 2017, pp. 136–145.
- [5] F. Facchinei and J.-S. Pang, *Finite-dimensional variational inequalities and complementarity problems*. Springer, 2003.